

# Lab 1 : Kubernetes cluster

---

## Step 1: Prerequisites Setup

On node1

```
# yum install vim -y
# hostname master
# echo 127.0.0.1 master >> /etc/hosts
```

logout and login to check the hostname setup on the master machine

On node2

```
# yum install vim -y
# hostname worker1
# echo 127.0.0.1 worker1 >> /etc/hosts
```

logout and login to check the hostname setup on the worker1 machine

On node3

```
# yum install vim -y
# hostname worker2
# echo 127.0.0.1 worker2 >> /etc/hosts
```

logout and login to check the hostname setup on the worker2 machine

---

## Step 2: Install Docker CE

Set up the repository and Install required packages.

```
# yum install yum-utils device-mapper-persistent-data lvm2
# yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo
```

Install Docker CE.

```
# yum update && yum install docker-ce-18.06.2.ce
```

# Restart Docker

```
systemctl daemon-reload
systemctl restart docker
systemctl enable docker
```

## Step 3: Install kubeadm

You will install these packages on all of your machines:

- kubeadm: the command to bootstrap the cluster.
- kubelet: the component that runs on all of the machines in your cluster and does things like starting pods and containers.
- kubectl: the command line util to talk to your cluster.

```
# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
```

```
# yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

```
# systemctl enable --now kubelet
```

#### **Step 4: Enable Net packet forwarding**

Some users on RHEL/CentOS 7 have reported issues with traffic being routed incorrectly due to iptables being bypassed. You should ensure net.bridge.bridge-nf-call-iptables is set to 1 in your sysctl config, e.g.

```
# cat <<EOF > /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
EOF
```

```
# sysctl --system
```

#### **Step 5: Setup master node**

The control-plane node is the machine where the control plane components run, including etcd (the cluster database) and the API server (which the kubectl CLI communicates with).

```
# kubeadm init --pod-network-cidr=192.16.0.0/16
```

To make kubectl work for your non-root user, run these commands, which are also part of the kubeadm init output:

```
# mkdir -p $HOME/.kube
# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

#### **Step 6: Enable Network CNI**

We are using calico to setup the pod cluster network

```
# kubectl apply -f https://docs.projectcalico.org/v3.8/manifests/calico.yaml
```

### **Step 7: Adding worker node**

Run the following command on the master node to generate command to add worker on the master node.

```
# kubeadm token create --print-join-command  
copy the output
```

Now login into each worker node and paste the command.

Once done, login back to the master node and run the following command to check the cluster status

```
# kubectl get nodes
```

You should be able to see all three nodes in ready state.

Congrats you have successfully deployed the kubernetes cluster using kubeadm.

## Lab 2: Pods lab

---

### Step 1: Namespace

Create a name space with the name new-ns with the manifest file

```
# vim ns.yml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: new-ns
```

save and quit the file

```
# kubectl create -f ns.yml
```

```
# kubectl get ns
```

### Step 2: Pod

Create a pod name pod-data with the manifest file in the new-ns namespace with nginx image

```
# vim pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod1
```

```
  namespace: new-ns
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

save and quit the file

```
# kubectl create -f pod.yml
```

```
# kubectl get pods -n new-ns
```

### Step 3: Pod with label

Create a pod named pod2 with label env=test with redis image

```
# vim pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod2
```

```
  labels:
```

```
    env: test
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: redis
```

save and quit the file

```
# kubectl create -f pod2.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl get pods --show-labels
```

## Lab 3: Replica's

---

### Step 1: RC

Setup a replica controller rc1 with 3 replicas and the pod selector to be env=prod

```
# vim rc.yml
```

```
apiVersion: v1
```

```
kind: ReplicationController
```

```
metadata:
```

```
  name: rc1
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    env: prod
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        env: prod
```

```
    spec:
```

```
      containers:
```

```
      - name: cont1
```

```
        image: nginx
```

save and quit the file

```
# kubectl create -f rc.yml
```

```
# kubectl get rc
```

```
# kubectl get pods -o wide
```

## Step 2: ReplicaSet

Create a Replica set rs1 with 3 replicas and two pod selection as env = test or env = prod

```
# vim rs.yml
```

```
apiVersion: apps/v1
```

```
kind: ReplicaSet
```

```
metadata:
```

```
  name: rs1
```

```
spec:
```

```
  replicas: 3
```

```
  selector:
```

```
    matchExpressions:
```

```
      - key: env
```

```
        operator: In
```

```
        values:
```

```
          - prod
```

```
          - test
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        env: prod
```

```
    spec:
```

```
      containers:
```

```
        - name: cont1
```

```
          image: nginx
```

save and quit the file

```
# kubectl create -f rs.yml
```

```
# kubectl get rs
```

```
# kubectl get pods
```

### Step 3: DaemonSet

Create the daemonset and see it's operation

```
# vim ds.yaml
```

```
apiVersion: apps/v1
```

```
kind: DaemonSet
```

```
metadata:
```

```
  name: ds1
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      abc: xyz
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        abc: xyz
```

```
    spec:
```

```
      containers:
```

```
        - name: newcont
```

```
          image: nginx
```

save and quit the file

```
# kubectl create -f ds.yaml
```

```
# kubectl get daemonset
```

```
# kubectl get pods -o wide
```



# Lab 3: Services

---

## Step 1: ClusterIP Service

Create a pod with label new=old and nginx image and expose it with clusterIP service on port 8080

```
# vim l3pod.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: l3pod
  labels:
    new: old
spec:
  containers:
    - name: cont1
      image: nginx
```

Save and quit the file

```
# kubectl create -f l3pod.yml
```

```
# kubectl get pods -o wide
```

```
# vim cip.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: cipservice
spec:
  type: ClusterIP
  selector:
    new: old
  ports:
    - name: port1
      port: 8080
      TargetPort: 80
```

save and quit the file

```
# kubectl create -f cip.yml
```

```
# kubectl get svc
```

```
# curl <serviceipaddress>
```

### Step 1: NodePort Service

Create a pod with label old=new and nginx image and expose it with clusterIP service on port 8080 and node port 32080

```
# vim l3pody.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: l3pody
  labels:
    old: new
spec:
  containers:
    - name: cont1
      image: nginx
```

Save and quit the file

```
# kubectl create -f l3pod.yml
```

```
# kubectl get pods -o wide
```

```
# vim nip.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: cipservice
spec:
  type: NodePort
  selector:
    new: old
  ports:
    - name: port1
      port: 8080
      TargetPort: 80
      nodePort: 32080
```

save and quit the file

```
# kubectl create -f nip.yml
```

```
# kubectl get svc
```

Open your desktop browser and access <http://mastereip:32080>

# Lab 4: Storages

---

## Step 1: Host path Volume

Create a host path volume to mount /data directory into /vishal of the container

```
# vim l4pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod3
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      volumeMounts:
```

```
        - name: vol1
```

```
          mountPath: /vishal
```

```
  volumes:
```

```
    - name: vol1
```

```
      hostPath:
```

```
        path: /data
```

```
# kubectl create -f l4pod.yml
```

```
# kubectl get pods
```

# Lab 5: Injecting Data

---

## Step 1: Configmap

Create a config map with name qwe with key value pair to be v1=vishal

Export this variable in the pod l5pod with environment variable name as username

```
# kubectl create configmap qwe --from-literal=v1=vishal
```

```
# kubectl get configmap
```

```
# vim l5pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: l5pod
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      env:
```

```
        - name: username
```

```
          valueFrom:
```

```
            configMapKeyRef:
```

```
              name: qwe
```

```
              key: v1
```

save and quit the file

```
# kubectl create -f l5pod.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl exec -it pod l5pod
```

inside the pod # echo \$username

the output should be vishal

## Step 2: Secret

Create a config map with name sec1 with key value pair to be pass=redhat

Export this variable in the pod l5pod with environment variable name as password

```
# kubectl create secret generic sec1 --from-literal=pass=redhat
```

```
# kubectl get secret
```

```
# vim l5pod1.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: l5pod1
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

```
      env:
```

```
        - name: username
```

```
          valueFrom:
```

```
            secretKeyRef:
```

```
              name: sec1
```

```
              key: pass
```

save and quit the file

```
# kubectl create -f l5pod.yml
```

```
# kubectl get pods -o wide
```

```
# kubectl exec -it pod l5pod
```

inside the pod # echo \$password

the output should be redhat

# Lab 6: Deployment

---

## Step 1: Deployment

Deploy a deployment dp1 with redis image and 3 replicas.

```
# vim l6dp.yml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: dp1
spec:
  replicas: 3
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: cont1
          image: redis
```

save and quit the file

```
# kubectl create -f l6dp.yml
```

```
# kubectl get deployment
```

```
# kubectl get pods -o wide
```

# Lab 7: Authentication

---

## Step 1: User creation

Create a user on the master machine then go into its home directory to perform the remaining steps.

```
# useradd jean && cd /home/jean
```

## Step 2: Certification generations

Create the private key and certificate signing request

```
# openssl genrsa -out jean.key 2048
# openssl req -new -key jean.key -out jean.csr -subj "/CN=jean"
```

## Step 3: Sign the CSR

Sign the CSR with the Kubernetes CA. We have to use the CA cert and key which are normally in /etc/kubernetes/pki/. Our certificate will be valid for 500 days

```
# openssl x509 -req -in jean.csr -CA /etc/kubernetes/pki/ca.crt -CAkey /etc/kubernetes/pki/ca.key -CAcreateserial -out jean.crt -days 500
```

## Step 4: Securing identification

Create a ".certs" directory where we are going to store the user public and private key.

```
# mkdir .certs && mv jean.crt jean.key .certs
```

## Step 5: K8s User creation

Create the user inside kubernetes

```
# kubectl config set-credentials jean --client-certificate=/home/jean/.certs/jean.crt
--client-key=/home/jean/.certs/jean.key
```

## Step 6: Context

Create a context for the user.

```
# kubectl config set-context jean-context --cluster=kubernetes --user=jean
```

## Step 7: Set RBAC

```
# mkdir /home/jean/.kube/config
```

*Ask the instruction from trainer to create the config file.*

```
# chown -R jean: /home/jean
```

```
# kubectl create ns mynctest
```

From root, create a file name rb.yml and put the following entry

```
# vim rb.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
```

```
metadata:
  name: jean
  namespace: mynctest
subjects:
- kind: User
  name: jean
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: ClusterRole
  name: edit
  apiGroup: rbac.authorization.k8s.io
```

```
# kubectl create -f rb.yml
# kubectl get rolebinding -n mynctest
```

login from jean user account and try creating the pod in mynctest namespace, you should be able to do the same.



# Lab 8: Security

---

## Scenario : Setting up host namespace in pod

Create a pod with network namespace to be shared with that of the host namespace.

Step1: Create pod manifest

Create the following manifest definition with details

```
# vim podns.yml
apiVersion: v1
kind: Pod

metadata:
  name: podns1

spec:
  hostNetwork: true
  containers:
    - name: cont1
      image: nginx
```

save and quit the file

```
# kubectl create -f podns.yml
# kubectl get pods -o wide
```

check the ip associate with the pod, it should be the ip address of the worker node on which it is deployed.

## Scenario : Using Security Context

In the scenario we will try to run a pod with root access denied as let's see if it run or throws the error.

Step 1: Create the manifest file with following details

```
# vim podsc.yml
apiVersion: v1
kind: Pod

metadata:
  name: podsc1

spec:
  containers:
    - name: cont1
      image: nginx
      securityContext:
        runAsNonRoot: true
```

Save and quit the file.

```
# kubectl create -f podsc.yml
```

```
# kubectl get pods -o wide
```

You should not be able to see the pod in the running state as there is some issue with the container configuration

```
# kubectl describe pod podsc1
```

Scroll down to the bottom to see the error mentioning the container needs root privilege to run.

### **Scenario : Managing Network Policies**

In this scenario, we will manage network policies

First scenario will be restricted pod to communicate with each other inside a namespace.

Step 1: Create a namespace and deploy two nginx pods inside it.

```
# vim pod.yml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: pod1
```

```
spec:
```

```
  containers:
```

```
    - name: cont1
```

```
      image: nginx
```

save and quit.

```
# kubectl create -f pod.yml
```

Repeat the same step by changing pod name to pod2 now.

Check the ip for both pods

```
# kubectl get pods -o wide
```

Login into pod1 and try doing the curl for another pod

```
# kubectl exec -it pod1 /bin/bash
```

*Inside the pod # curl <ip of second pod>*

You should be able to connect with the nginx application running on the pod.

Step 2: Apply the default deny network policy

```
# vim default-ns.yml
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: default-deny
```

```
spec:
```

```
  podSelector: {}
```

save and quit the file

```
# kubectl create -f default-ns.yml
```

```
# kubectl get networkpolicy
```

Now login into pod1 again and try to access nginx on pod2

```
# kubectl exec -it pod1 /bin/bash
```

inside the pod # curl ip of pod 2

This time you should not be able to connect.

Step 3: Allow controlled traffic between these two pods

```
# kubectl label pod pod1 env=test
```

```
# kubectl label pod pod2 env=prod
```

write the network policy to enable the communication from pod1 to pod2 for port 80.

```
# vim ns.yml
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: allow-nspolicy
```

```
spec:
```

```
  podSelector:
```

```
    matchLabels:
```

```
      env: prod
```

```
  ingress:
```

```
    - from:
```

```
      - podSelector:
```

```
        matchLabels:
```

```
          env: test
```

```
        ports:
```

```
    - port: 80
```

save and quit the file.

```
# kubectl create -f ns.yml
```

Follow the steps to login into pod1 and try access pod2

You should be able to access it this time.

# Lab 8: Computing Resources

---

## Scenario : Set request

In this exercise we will setup the request for container in a pod definition

Step 1: Create the pod resource with following resource request

```
# vim podcr.yml
apiVersion: v1
kind: Pod
metadata:
  name: podcr1

spec:
  containers:
    - image: busybox
      command: ["dd", "if=/dev/zero", "of=/dev/null"]
      name: main
      resources:
        requests:
          cpu: 200m
          memory: 10Mi
```

save and quit the file.

```
# kubectl create -f podcr.yml
```

To verify the memory allocated run the following command to check the status

```
# kubectl exec -it podcr1 top
```

In the output you should be able to see the memory allocated at the top.

## Scenario : Create limit range

In this scenario, we will be creating the limit range to define default limit and request for a namespace.

```
# kubectl create ns limit-ns
```

```
# vim limitns.yml
apiVersion: v1
kind: LimitRange
metadata:
  name: example
spec:
  limits:
    - type: Pod
      min:
        cpu: 50m
        memory: 5Mi
      max:
```

```
    cpu: 1
    memory: 1Gi
- type: Container
defaultRequest:
  cpu: 100m
  memory: 10Mi
default:
  cpu: 200m
  memory: 100Mi
min:
  cpu: 50m
  memory: 5Mi
max:
  cpu: 1
  memory: 1G
```

save and quit the file

```
# kubectl create -f limitns.yml -n limit-ns
# kubectl get limitrange -n limit-ns
```

### **Scenario : Using Node affinity**

In this scenario, we will deploy the pod to a node with nodeselector option

```
# kubectl label node worker1 disk=ssd
# kubectl get nodes --show-labels
```

```
# vim podsel.yml
apiVersion: v1
kind: Pod
```

```
metadata:
  name: podsel1
```

```
spec:
  nodeSelector:
    disk: ssd
  containers:
- name: cont1
  image: nginx
```

save and quit the file.

```
# kubectl creat -f podsel.yml
# kubectl get pods -o wide
```

The output should show you that the pod is deployed on the worker node 1